

Week 1 deliverable — three questions

Seat	team06@theschoolofai.in · roles sales_user, user, agent_user · apps crm, agent
Evidence	tools/list → 247 tools / 51 entities · live SalesOrder.list (total: 310, limit: 20) · rechecked against tools-list.json 2026-09-22
Compared against	Salesforce Agentforce · HubSpot Breeze · Attio · Zoho CRM · ERPNext
Full register	gap-analysis.md · AgentSwitch-Gap-Report.docx

One page. Three questions. Specifics only. Tool names below are exact catalog strings.

1. What do they do that we do not?

Concrete absences in the CRM/agent surface this seat can see — not vibes.

Gap	Evidence in our MCP surface	Who ships it
No aggregate / group-by reporting	No *.aggregate tool in the 247. SalesOrder.list returns 20 of 310; "revenue this quarter by owner" is 16 round trips plus client-side sum	Zoho reports, Agentforce + Data Cloud, ERPNext
No CRM invoice, payment, or credit note	Quote → order exists (Quotation.make.SalesOrder, also Quotation.convert_to_order.*). Live orders sit at invoiced_status: "unbilled". No Invoice.* / CreditNote.* / SalesOrder.make.Invoice. (Storefront has endpoint.storefront.payment.verify — commerce checkout, not AR billing.)	Zoho, Odoo, ERPNext
GST fields with no filing path	Order carries gst_treatment, place_of_supply, gst_no, tds_*, tcs_*; line items carry hsn_or_sac. No IRN, QR, or GSTR export tool	ERPNext native e-invoice
Agents can be run, not authored	AgentSkill, AgentPersona, AgentRunbook, AgentToolPolicy, AgentProvider, AgentFloorConfig, AgentJob are list/get only (AgentRunbookRun likewise). AgentPersona.daily_limits is a separate read-only route	Agentforce Builder, Breeze Studio
No knowledge grounding / RAG	FileAttachment is create/list/get/update of blobs; no Knowledge.ingest / Knowledge.search	Agentforce Intelligent Context, Copilot Studio
No service desk or SLA clock	BugReport: create / list / get only — no update, no queue, no SLA, no CSAT. No Ticket.*	HubSpot Service Hub, Zoho Desk
No outbound email, sequences, or templates	Sole email tool: endpoint.email.public.subscriber_lists (read). No Email.send, no Sequence.*	Attio Pro, HubSpot, Zoho
No forecast, quota, or territory rollup	Goal, Pipeline, AccountPlan are list/get only. Deal.create.probability is a typed	Zoho Enterprise, Clari

	number (default 0), not a model output	
No dedupe, merge, or enrichment	Party : create/list/get/update. No Party.find_duplicates , Party.merge , or enrichment resolve	Clay, Breeze Intelligence, Attio
Interactions name only one person	Note.create and Activity.create each take a single party_id (plus optional deal_id / lead_id / company_id). No participant array	Every calendar-integrated CRM
Deals are single-threaded	Deal.create takes products as an array but contact_id and party_id as scalars . source enum includes referral with no referrer party field	Salesforce opportunity contact roles, HubSpot deal contacts
Relationship edges exist but are inert	PartyRelationship has create/list/get/update over 16 typed edges (guardian ... counterparty , including referred_by / referred). Nothing observes, proposes, or walks them — and there is no delete	Nobody holds <i>observed</i> graph; TeamLink is outside CRM

2. Which of those can an agent close with tools this seat already has?

Two buckets. Do not confuse them.

Ours to fix (new tables / endpoints — agent cannot invent these)

Missing capability	Why the seat cannot close it by orchestration
{Entity}.aggregate	No group-by tool exists; paginating 310 orders and summing in the model is not reporting
Invoice / CreditNote + SalesOrder.make.Invoice	Revenue chain dead-ends at SalesOrder; GST filing has nowhere to hang. Storefront payment.verify does not bill a sales order
AgentSkill.create / AgentPersona.create / version / publish	Governance entities are read-only; an agent cannot author what it cannot write
Knowledge.ingest / Knowledge.search	Attachments are unindexed; no retrieval surface
Ticket.* with SLA	BugReport has no update path and no queue semantics
Email.send / Sequence.*	No send tool on any channel despite AgentSession.create.channel ∈ {web, whatsapp, telegram, slack, discord, email, api}
ActivityParticipant join	Schema forbids multi-party interactions; co-attendance cannot be expressed
DealContact + referred_by_party_id	Buying committee and referral attribution are not columns an agent can invent

Yours to build (orchestration over what already exists)

Agent goal	Tools this seat already has (exact names)	What the agent must do
Move a lead → deal → quote → order	Lead.qualify.* / Lead.make.Deal → Deal.qualify → Deal.send_proposal → Deal.negotiate → Deal.mark_won.* / Deal.mark_lost.* → Deal.make.Quotation → Quotation.send /	Hold the stage goal, re- get after each transition, branch on failure

	<code>Quotation.accept.* → Quotation.make.SalesOrder</code> (or <code>Quotation.convert_to_order.*</code>) → <code>SalesOrder.confirm / SalesOrder.approval.submit</code>	
File notes and activities on a party/deal	<code>Note.create / .update; Activity.create / .update; CallNoteDraft.list / .get</code> (read-only drafts — no approve/publish tool)	Draft in <code>CallNoteDraft</code> if present, then write Note / Activity yourself; escalate if human edit is required
Respect spend and reach caps	<code>AgentPersona.daily_limits</code> (tokens/USD used vs budget, blocked flag, reset 00:00 UTC); <code>AgentToolPolicy.list / .get; AgentFloorConfig.list / .get</code>	Refuse or escalate when limits are exhausted before calling spendy tools
Escalate when stuck	<code>AgentEscalation.create / .list / .get</code>	Park the goal, hand a human a decision, resume after resolution
Prove what the agent did	<code>endpoint.job_ledger.verify / endpoint.job_ledger.replay / endpoint.job_ledger.forensics; AgentLedgerSeal.list / .get</code>	After a multi-step run, produce a replayable audit trail
Honor DSAR / retention	<code>PrivacyRequest.create / .update; endpoint.agent_governance.privacy.forget.preview; AgentRetentionRule.list; LedgerRetentionPolicy.list; endpoint.job_ledger.retention.preview</code>	Preview blast radius, then execute forget under policy
Operate the agent floor	<code>endpoint.mission_control.staffing_forecast / .incidents / .incident_timeline / .runbooks; AgentRunbook.list / AgentRunbookRun.list</code>	Treat agents as a staffed floor, not toggles
Keep working memory across steps	<code>AgentMemory.create / .list; AgentTodo.create / .list; AgentTask.create / .list</code>	Persist intermediate state so a twenty-step goal survives session boundaries
Manual relationship hygiene	<code>PartyRelationship.create / .update / .list / .get; Party.list / .get</code>	Type declared edges (referred_by, employer, billing_contact, ...) — not observed inference

If a gap needs a column that is not there, stop and file it under *ours*. If every step already has a tool, that is *yours* — write the skill / runbook, do not wait for platform.

3. What can an agent do that their product cannot?

Their product is a UI over a database: a human clicks every step. Our seat is a **permissioned tool surface plus a job ledger**. That changes what “using the CRM” means.

Their UI forces	An AgentSwitch agent can
One screen, one action, one human	Hold a goal across twenty tool calls — lead qualify → deal negotiate → quotation → sales order confirm — re-reading <code>Deal.get / Quotation.get / SalesOrder.get</code> when state changes underneath
Telemetry dashboards after the fact	Prove a specific action with <code>endpoint.job_ledger.replay / .forensics</code> and <code>AgentLedgerSeal</code> — cryptographic answer to “ <i>what did the agent do and why</i> ”
Soft policy documents	Read <code>AgentToolPolicy + AgentPersona.daily_limits + AgentFloorConfig</code> before execution; hard-stop or escalate when blocked

Client-side “are you sure?” on writes	Server-side caller-scoped discovery (<code>_meta.agentswitch.discovery: "caller"</code>) — <code>tools/list</code> only returns what this seat’s roles allow
Privacy as a ticket queue	Runtime DSAR: <code>PrivacyRequest</code> + <code>endpoint.agent_governance.privacy.forget.preview</code> with blast radius before delete
Pipeline fields a human typed (<code>probability</code> , stage)	(Opening, not yet built) Walk an observed relationship graph at deal time — buying committee, warm path, champion decay — with every edge traceable via the ledger. Schema seed exists (<code>PartyRelationship</code> , 16 types); capture and traversal are still ours to ship

The line to hold in a buyer conversation: competitors sell *pipeline* intelligence (attributes of a record). We are built to sell *agent* intelligence (a goal held across mutating state, under policy, with proof). Relationship intelligence on deals is the product differentiator to build next; governance + ledger endpoints are the differentiators already shipped and under-marketed.

Rechecked 2026-09-22 against `tools-list.json` (247 tools). Corrections from prior draft: ledger/mission-control tools are `endpoint.job\_ledger` / `endpoint.mission\_control`; forget preview is `endpoint.agent\_governance.privacy.forget.preview`; `AgentPersona.daily\_limits` is its own tool; `CallNoteDraft` has no approve path; `hsn\_or\_sac` lives on line items; storefront `payment.verify` is not CRM AR.